

Retrieval-Augmented Distributed Image Compression with Multi-Reference Side Information

Guojun Xu[✉], Jianwen Xiang[✉], *Member, IEEE*, Mingyang Zhang[✉], Yaning Xie,
Junwei Zhou[✉], *Member, IEEE*

Abstract—Extremely low-bitrate distributed image compression (DIC) is crucial for bandwidth-constrained visual sensing and distributed cloud storage. At such low bitrates (<0.1 bpp), the transmitted representation retains little target information, making reconstruction particularly sensitive to the relevance and geometric consistency of the available side information (SI). Existing learning-based DIC methods are commonly developed with a pre-associated SI under matched training and evaluation distributions, whereas practical decoder-side visual collections may contain references with varying correlations to the target. To address these limitations, we propose RADIC, a retrieval-augmented DIC framework that replaces fixed paired SI with multiple references retrieved from a decoder-side visual gallery. Specifically, we construct a heterogeneous SI gallery and retrieve correlated references using the transmitted latent representation. To handle the geometric mismatch and varying relevance of retrieved references, we design a depth-aware geometric alignment module to align retrieved features through depth-guided coarse flow and residual refinement. A multi-reference cross-attention mechanism then selectively aggregates complementary information while suppressing irrelevant content. Moreover, decoder attention supervises an encoder-side importance map for SI-aware rate allocation. Trained exclusively on KITTI General, RADIC achieves superior perceptual compression performance across KITTI General, KITTI Stereo, Cityscapes, and InStereo2K, and remains robust under weakly correlated and cross-domain SI.

Index Terms—Distributed image compression, side information, image retrieval, perceptual compression, Extremely low bitrates.

I. INTRODUCTION

LARGE-SCALE multi-view systems, such as autonomous driving [1], robotic perception [2], and immersive communication [3], continuously generate correlated visual observations from different viewpoints. Efficient transmission of these observations is challenging when each view is captured and encoded independently under limited bandwidth. Meanwhile, practical

This work was partially supported by the Hubei Province Major Science and Technology Innovation Program (2024BAA011), the Key Research and Development Program of Hubei Province (2025BEB012).

Guojun Xu, Jianwen Xiang, Mingyang Zhang, Yaning Xie, and Junwei Zhou were with the School of Artificial Intelligence, Wuhan University of Technology, Wuhan, China, 430070. Corresponding author: Junwei Zhou (E-mail: junweizhou@msn.com).




Fig. 1: Comparison of two DIC settings. (a) Existing methods [6], [7] rely on a paired SI for each target. (b) The proposed RADIC retrieves top- K references from a decoder-only visual database, aligns them geometrically, and uses the aligned features for reconstruction, thereby relaxing the paired-SI assumption.

decoders, such as cloud servers and edge storage nodes, may maintain large collections of previously stored visual data that can provide additional information for reconstruction [4]. This setting motivates distributed image compression (DIC), where the target image is encoded independently and side information (SI) available at the decoder is used to assist reconstruction [5].

As illustrated in Fig. 1, existing learning-based DIC methods are generally developed under a paired-SI setting, where each target is associated with a pre-matched stereo or multi-view reference [6]–[10]. In practical systems, however, the available SI may exhibit different levels of correlation with the target. As shown in Fig. 2, the same target can produce substantially different reconstructions when paired, non-paired, and cross-domain SI are provided to the decoder. As the correlation between target and SI decreases, larger differences in geometry and image content are introduced, making the SI less effective for reconstruction. In addition, target images may be acquired from different sensors and environments, resulting in changes in the source distribution. These effects become more significant at extremely low bitrates, where only limited target information is retained in the transmitted representation.




Fig. 2: Qualitative comparison of ATN, NDSC, MDIC, and RADIC under different decoder-only SI conditions. The top row shows the target image and the cross-domain ($\triangle$), non-paired ($\diamond$), and paired ($\circ$) SI. For each method, the three reconstruction rows correspond to these SI conditions from top to bottom. The values below each reconstruction report Bpp/ PSNR/ LPIPS.

At extremely low bitrates, perceptual reconstruction becomes important because pixel-wise optimization tends to produce over-smoothed results when image details are largely removed from the transmitted representation [11]. Recent single-image compression methods [12], [13] have introduced generative priors to recover perceptually plausible structures and textures. MDIC [7] further extends perceptual reconstruction to DIC by exploiting multimodal SI. However, its visual SI is still provided as a pre-matched reference, and the reconstruction remains sensitive to the correlation between the target and the given SI.

To address these challenges, as illustrated in Fig. 1, we propose RADIC, a retrieval-augmented DIC framework for perceptual reconstruction at extremely low bitrates. At the decoder, the decoded target representation is used to retrieve the top- K candidate references from a visual gallery containing SI from multiple domains. By decoupling SI acquisition from domain-specific pairing, RADIC preserves target-only encoding while adapting decoder-side reconstruction to different visual domains and SI correlations.

Retrieval from a heterogeneous gallery introduces substantial geometric mismatch and relevance variation among the candidate references. To establish the corre-

spondence between target and references, we introduce a Depth-Aware Geometric Alignment (DAGA) module that combines depth-guided coarse flow with residual refinement. A Multi-reference Cross-Attention (MCA) mechanism then performs spatially selective aggregation, retaining complementary information while suppressing irrelevant content. To improve robustness to varying SI quality, decoder attention further supervises an encoder-side importance map during training, transferring SI-aware guidance to rate allocation. As shown in Fig. 2, RADIC preserves clearer scene structures under paired, non-paired, and cross-domain SI, showing improved robustness to changes in SI correlation.

Our main contributions are summarized as follows:

- We formulate a retrieval-augmented DIC setting in which multiple SI references are selected from a decoder-only visual gallery according to the received target representation, enabling perceptual reconstruction at extremely low bitrates under varying correlations between target and references.
- We address the geometric mismatch and varying relevance of retrieved SI through depth-aware alignment and selective multi-reference fusion, which exploit complementary information while suppressing irrelevant content. Decoder attention further guides

importance learning for SI-aware rate allocation.

- RADIC is trained only on KITTI General and evaluated on KITTI General, KITTI Stereo, Cityscapes, and InStereo2K. Experiments show favorable perceptual performance across four test sets under the stated shared-gallery protocol.

The remainder of this paper is organized as follows. Section II reviews related work on correlated-source image compression and extremely low-bitrate perceptual compression. Section IV presents the proposed RADIC framework. Section V describes the experimental settings and evaluates the proposed method in terms of compression performance and robustness. Section VI further analyzes the use of multiple SI references under different SI conditions. Finally, Section VII concludes the paper.

II. RELATED WORK

A. Correlated-Source Image Compression

Correlated visual sources can be compressed through two main paradigms. Joint image compression (JIC) accesses correlated images at both the encoder and decoder, allowing direct inter-source interaction and joint entropy modeling. ECSIC [14] introduces epipolar-constrained stereo cross-attention and stereo context modules for entropy estimation. BiSIC [15] performs bidirectional stereo compression with a cross-dimensional entropy model to capture dependencies between stereo latents. CAMSIC [16] employs content-aware masked image modeling to estimate spatial and disparity dependencies. 3D-LMVIC [17] further incorporates 3D geometric priors for wide-baseline multi-view compression. Although these methods effectively exploit inter-source correlations, they rely on encoder-side interaction and are unsuitable for distributed scenarios.

Distributed image compression (DIC) encodes correlated sources independently and exploits their dependencies using side information (SI). Slepian–Wolf theory [18] shows that lossless distributed coding can achieve the joint-coding rate region, while Wyner–Ziv coding [19] extends this setting to lossy compression. In learning-based DIC, NDIC [5] extracts common information from SI and combines it with the transmitted target representation. ATN [10] uses cross-attention to align target and SI features. NDSC [6] formulates lossy distributed coding with a conditional VQ-VAE, while LDMIC [9] employs symmetric DIC for multi-view reconstruction. More recently, DIC-DDA [20] improves asymmetric DIC through dual-domain alignment between the decoded target and SI.

These methods achieve competitive rate-distortion performance on stereo and multi-view datasets, but most are optimized for pixel fidelity and produce blurred reconstructions at extremely low bitrates. They also commonly rely on a single paired SI image and in-domain

training and evaluation. Consequently, their performance is sensitive to changes in SI correlation and limited in generalizing across datasets.

B. Extremely Low-Bitrate Perceptual Compression

At extremely low bitrates, perceptual optimization helps preserve richer semantic information and alleviate the blurring caused by severe information loss [13], [21], [22]. Early perceptual compression methods employ generative models to improve reconstruction quality at low rates [23]. More recent approaches make use of stronger semantic and diffusion priors. PerCo [11] conditions an iterative diffusion decoder on vector-quantized latents and a global description, while ResULIC [24] combines residual-guided semantic coding with compression-aware diffusion. MISC [25] introduces large multimodal models for ultra-low bitrate semantic compression. These methods focus on single-image compression and do not exploit decoder-only correlated sources.

Perceptual reconstruction has also been introduced into DIC. MDIC [7] exploits multimodal decoder-side SI for extremely low-bitrate reconstruction, while CODIC [26] considers distortion-perception optimization in distributed compression. However, both methods are developed with pre-associated SI and do not explicitly consider large variations in target–SI correlation. Their reconstruction remains dependent on the quality of the given SI. We address these limitations by retrieving multiple references from a heterogeneous multi-domain SI gallery and selectively exploiting correlated information while suppressing irrelevant content.

III. INFORMATION-THEORETIC ANALYSIS

A. Problem Formulation

Let $X^n = (X_1, \dots, X_n)$ denote a target source block of length n , taking values in $\mathcal{X}^n$. Let $Y_{1:K}^n = (Y_1^n, \dots, Y_K^n)$ denote K SI sources available only at the decoder. The encoder maps the target block to a message

$$f_n : \mathcal{X}^n \rightarrow \{1, \dots, M_n\}, \quad (1)$$

where M_n is the number of possible messages and satisfies

$$\frac{1}{n} \log_2 M_n \leq R. \quad (2)$$

For an observed source block X^n , the transmitted message is

$$M = f_n(X^n). \quad (3)$$

Here, R is the average number of transmitted bits per source symbol. Equivalently, one may set $M_n = \lfloor 2^{nR} \rfloor$.

The decoder reconstructs the target using the message and SI:

$$\hat{X}^n = g_n(M, Y_1^n, \dots, Y_K^n). \quad (4)$$




Fig. 3: Overall architecture of RADIC. The encoder maps the target image I_x to a latent representation z_x and predicts an SI-aware importance map M_{pred} for latent modulation before quantization. At the decoder, $\hat{z}_x$ retrieves top- K references from the gallery $\mathcal{G}$, which are aligned by Depth-Aware Geometric Alignment (DAGA) and screened by Multi-reference Cross-Attention (MCA). The fused reference features condition a latent diffusion model to reconstruct $\hat{I}_x$ at extremely low bitrates.

For a single-letter distortion measure $d(x, \hat{x})$, an (n, R, D) code satisfies

$$\frac{1}{n} \sum_{i=1}^n \mathbb{E} [d(X_i, \hat{X}_i)] \leq D. \quad (5)$$

B. Condition for Multi-Reference Gain

For lossy reconstruction, the Wyner–Ziv rate–distortion function with K decoder-only references is

$$R_{\text{WZ}}^{(K)}(D) = \min_{p(u|x), \hat{x}(u, y_{1:K})} I(X; U | Y_{1:K}), \quad (6)$$

$$\mathbb{E}[d(X, \hat{X})] \leq D$$

where U denotes the compressed representation generated from X , and

$$\hat{X} = \hat{x}(U, Y_{1:K}). \quad (7)$$

The Markov chain

$$U - X - Y_{1:K} \quad (8)$$

means that U and $Y_{1:K}$ are conditionally independent given X :

$$p(u | x, y_{1:K}) = p(u | x). \quad (9)$$

This enforces the decoder-only SI constraint because the encoder generates U without observing any reference.

Define the optimal rate gain from adding the $(K+1)$ -th reference as

$$\Delta_{K+1}(D) = R_{\text{WZ}}^{(K)}(D) - R_{\text{WZ}}^{(K+1)}(D). \quad (10)$$

The additional reference strictly improves compression efficiency if and only if

$$\Delta_{K+1}(D) > 0. \quad (11)$$

Equivalently, assuming that the minimum is attained, there must exist a feasible representation and reconstruction function for the $(K+1)$ -reference problem such that

$$I(X; U | Y_{1:K}, Y_{K+1}) < R_{\text{WZ}}^{(K)}(D). \quad (12)$$

For a fixed representation U satisfying

$$U - X - (Y_{1:K}, Y_{K+1}), \quad (13)$$

the reduction in conditional mutual information is

$$I(X; U | Y_{1:K}) - I(X; U | Y_{1:K}, Y_{K+1}) \\ = I(U; Y_{K+1} | Y_{1:K}). \quad (14)$$

For the current compressed representation, the $(K+1)$ -th reference yields a strict rate gain if and only if

$$I(U; Y_{K+1} | Y_{1:K}) > 0. \quad (15)$$

Eq. (15) means that Y_{K+1} must provide information about U that is not already contained in the first K references. The new reference is conditionally redundant and cannot further reduce the required rate for that representation when

$$I(U; Y_{K+1} | Y_{1:K}) = 0, \quad (16)$$

Hence, increasing K is beneficial only while newly added references provide complementary target information. As their conditional information decreases, the rate gain approaches zero and eventually saturates.

IV. METHODOLOGY

A. Overall Framework

The proposed RADIC framework consists of an independent encoder and a retrieval-augmented decoder, as illustrated in Fig. 3. Given a target image $I_x \in \mathbb{R}^{H \times W \times 3}$, a VAE encoder $\mathcal{E}_{vae}$ first maps it into a latent representation, which is further processed by an analysis transform g_a to obtain the compression latent z_x^0 . Before quantization, an SI-aware Importance Map (IM) modulates z_x^0 according to the predicted availability of decoder-only SI. The resulting latent is quantized and entropy-coded into the transmitted bitstream.

At the decoder, the decoded latent $\hat{z}_x$ is used to retrieve the top- K candidate references from a decoder-only visual gallery $\mathcal{G}$. The retrieved references are aligned with the target representation by the Depth-Aware Geometric Alignment (DAGA) module. A Multi-reference Cross-Attention (MCA) module then performs spatial reference screening and feature fusion. The fused reference features, together with the decoded target representation, condition a latent diffusion model for reconstruction at extremely low bitrates.

B. Encoding and SI-Aware Rate Allocation

In the distributed coding setting, the encoder has no access to SI. It cannot determine which target regions can be recovered from the available SI. RADIC addresses this issue by learning an encoder-side importance map from the target representation and supervising it with decoder-side attention scores during training.

The target image is first mapped to the compression latent

$$z_x^0 = g_a(\mathcal{E}_{vae}(I_x)). \quad (17)$$

A mask predictor $\mathcal{P}_m$ then estimates a spatial IM

$$M_{pred} = \mathcal{P}_m(z_x^0), \quad (18)$$

where $M_{pred} \in [0, 1]^{H_z \times W_z}$, and the modulated latent is

$$z_x = z_x^0 \odot (1 - \eta M_{pred}), \quad (19)$$

where $\odot$ denotes element-wise multiplication and η controls the attenuation strength.

C. Quantization and Entropy Modeling

Following RDEIC [13], a hybrid quantization scheme is adopted. The modulated latent z_x is quantized by rounding:

$$\hat{z}_x = Q(z_x). \quad (20)$$

A hyper-analysis transform h_a produces the hyper-latent

$$z_{h,x} = h_a(z_x), \quad (21)$$

which is quantized using a VQ codebook [27] to obtain $\hat{z}_{h,x}$.

For entropy coding, $\hat{z}_x$ is divided into ten channel slices, and each slice is partitioned into checkerboard anchor and non-anchor symbols. For a symbol $\hat{z}_{x,s,u}$ at position u in slice s , its Gaussian likelihood is conditioned on the quantized hyper-latent, the previously decoded slices, and the available checkerboard context. The corresponding rate is estimated as

$$R_{z_x} = -\frac{1}{HW} \sum_s \sum_u \log_2 p_{\psi_{s,u}}(\hat{z}_{x,s,u} \mid \hat{z}_{h,x}, \hat{z}_{x,<s}, \mathcal{C}_{s,u}), \quad (22)$$

where $\mathcal{C}_{s,u}$ denotes the available spatial context. Anchor symbols are decoded first and are subsequently used as context for non-anchor symbols.

The hyper-latent is quantized using a codebook $\mathcal{E}$ with $|\mathcal{E}| = 2^{14}$. Since each VQ index corresponds to a 64×64 region of the input image, its fixed rate contribution is

$$R_{z_h} = \frac{H_h W_h \log_2 |\mathcal{E}|}{HW} = \frac{14}{64^2} \simeq 0.003418 \text{ bpp}. \quad (23)$$

The bitrate is

$$\mathcal{R} = R_{z_x} + R_{z_h}. \quad (24)$$

The VQ branch is optimized with

$$\mathcal{L}_{VQ} = \|\text{sg}[z_{h,x}] - e\|_2^2 + \lambda_{\text{com}} \|z_{h,x} - \text{sg}[e]\|_2^2, \quad (25)$$

where e denotes the selected codebook embedding, λ_{com} controls the commitment term, and $\text{sg}[\cdot]$ denotes the stop-gradient operator.

D. Retrieval-Augmented Generative Decoding

Upon receiving the quantized latent $\hat{z}_x$, the decoder first reconstructs the target feature $f_{\text{main}} = g_s(\hat{z}_x)$, which is further compressed along the channel dimension to obtain the latent representation $\hat{z}_{Dx} = \mathcal{C}_c(f_{\text{main}})$. Consistent with the target encoding process, for each retrieved side image I_{y_k} ($k \in \{1, \dots, K\}$), we similarly extract $f_{\text{side}}^{(k)}$ and z_{Dy_k} .

1) *Depth-Aware Geometric Alignment (DAGA)*. Multiple SI sources inevitably suffer from complex geometric distortions. To align these features, we propose the DAGA module (Fig. 4) driven by monocular depth priors.




Fig. 4: Depth-Aware Geometric Alignment (DAGA). It predicts global affine parameters (α, β) to construct a depth-guided coarse flow, refines it with a residual flow, and warps the K side features. A learned confidence mask c_f is used to suppress unreliable regions.

For each side image, we extract its depth map D_{side} using the pre-trained Depth-Anything V3 [28]. The depth map is resized, inverted, and normalized to obtain the inverse depth d_{inv} . To model global viewpoint differences, a lightweight network predicts the affine parameters α and β from f_{main} and $f_{side}^{(k)}$. A coarse horizontal flow field is then generated by

$$flow_x^{coarse} = \alpha \cdot d_{inv} + \beta. \quad (26)$$

To refine the coarse alignment, a residual flow $\Delta flow = (\Delta flow_x, \Delta flow_y)$ and a confidence mask $c_f \in [0, 1]^{H \times W}$ are estimated from d_{inv} , $flow_x^{coarse}$, $f_{side}^{(k)}$, and f_{main} . The final flow field is defined as

$$D_X = flow_x^{coarse} + \Delta flow_x, \quad D_Y = \Delta flow_y. \quad (27)$$

The aligned side feature is then obtained by

$$\tilde{f}_{side}^{(k)} = \text{GridSample}(f_{side}^{(k)}, (D_X, D_Y)) \odot c_f, \quad (28)$$

where $\odot$ denotes pixel-wise multiplication.

2) *Multi-reference Cross-Attention (MCA)*. Given the target feature f_{main} and the K aligned reference features $\tilde{f}_{side} \in \mathbb{R}^{B \times K \times C \times H \times W}$, MCA selects one reference at each spatial location according to the similarity between target and reference. The target feature is projected into a query, while the aligned reference features are projected into keys and values:

$$Q = W_Q f_{main}, \quad K = W_K \tilde{f}_{side}, \quad V = W_V \tilde{f}_{side}. \quad (29)$$

For each spatial location (i, j) , the similarity score of the k -th reference is

$$s_k(i, j) = \frac{Q(i, j) \cdot K_k(i, j)}{\sqrt{C}}. \quad (30)$$

The hard selection mask is defined as

$$M_{Hard}^{(k)}(i, j) = \begin{cases} 1, & k = \arg \max_{k'} s_{k'}(i, j), \\ 0, & \text{otherwise.} \end{cases} \quad (31)$$

A soft mask is computed only for gradient propagation:

$$M_{Soft}^{(k)}(i, j) = \frac{\exp(s_k(i, j))}{\sum_{k'} \exp(s_{k'}(i, j))}. \quad (32)$$

The straight-through mask is

$$M_{STE}^{(k)} = M_{Soft}^{(k)} + \text{sg}(M_{Hard}^{(k)} - M_{Soft}^{(k)}), \quad (33)$$

Thus, the forward pass uses the one-hot hard mask, while gradients are propagated through $M_{Soft}^{(k)}$.

The selected reference feature is

$$f_{fused}(i, j) = \sum_{k=1}^K M_{STE}^{(k)}(i, j) \cdot V_k(i, j). \quad (34)$$

The same MCA operation is applied to the aligned latent representations:

$$z_{fused} = \text{MCA}(z_{Dx}, \{\tilde{z}_{Dy_k}\}_{k=1}^K). \quad (35)$$

3) *Feature and Latent Fusion*. The fused representations are injected into the decoding process through residual modulation:

$$\tilde{f}_{main} = f_{main} + \phi_f([f_{main}, f_{fused}]), \quad (36)$$

$$\tilde{z}_{Dx} = \phi_g([\hat{z}_{Dx}, z_{fused}]) \odot \hat{z}_{Dx} + \phi_z([\hat{z}_{Dx}, z_{fused}]), \quad (37)$$

where ϕ_f, ϕ_z denote learnable fusion networks and $\phi_g(\cdot)$ is a gating function.

4) *Decoding via Latent Diffusion Model*. Inspired by RDEIC [13], we adopt relay residual diffusion for latent reconstruction. Let $x_0 = \mathcal{E}_{vae}(I_x)$ denote the target VAE latent, and let $\tilde{z}_{Dx}$ be the SI-enhanced compressed latent. Their residual is

$$e_x = \tilde{z}_{Dx} - x_0. \quad (38)$$

Instead of starting from pure noise, the diffusion trajectory is initialized from the compressed latent:

$$x_N = \sqrt{\bar{\alpha}_N} \tilde{z}_{Dx} + \sqrt{1 - \bar{\alpha}_N} \epsilon_N, \quad \epsilon_N \sim \mathcal{N}(0, I). \quad (39)$$

To make the trajectory agree with the compressed-latent initialization at $n = N$, we use $\lambda = (\bar{\alpha}_N^{-1} - 1)^{-1/2} = \sqrt{\bar{\alpha}_N / (1 - \bar{\alpha}_N)}$. The relay residual diffusion process is defined as

$$x_n = \sqrt{\bar{\alpha}_n} x_0 + \sqrt{1 - \bar{\alpha}_n} (\lambda e_x + \epsilon_n), \quad (40)$$




Fig. 5: Rate-perception comparisons on four test datasets. Lower ($\downarrow$) is better for all metrics.

where the denoiser predicts the combined residual noise conditioned on the fused feature $\tilde{f}_{main}$:

$$\epsilon_{\theta}(x_n, n | \tilde{f}_{main}). \quad (41)$$

Starting from x_N , the reverse process estimates $\hat{x}_0$, which is decoded as

$$\hat{I}_x = \mathcal{D}_{\text{vae}}(\hat{x}_0). \quad (42)$$

E. Loss Functions

The overall training objective is defined as

$$\mathcal{L} = \lambda_1 \mathcal{L}_{diff} + \lambda_2 \mathcal{L}_{rate} + \lambda_3 \mathcal{L}_{guide} + \lambda_4 \mathcal{L}_{IM}. \quad (43)$$

Following the relay residual diffusion process, the diffusion loss supervises the prediction of the combined compression residual and Gaussian noise:

$$\mathcal{L}_{diff} = \omega_n \mathbb{E} \left[\left\| \lambda e_x + \epsilon_n - \epsilon_{\theta}(x_n, n | \tilde{f}_{main}) \right\|_2^2 \right], \quad (44)$$

where $e_x = \tilde{z}_{Dx} - x_0$ and ω_n is the timestep-dependent weighting coefficient.

The rate loss combines the estimated coding rate and the vector-quantization objective:

$$\mathcal{L}_{rate} = \mathcal{R} + \mathcal{L}_{VQ}. \quad (45)$$

The latent guidance loss constrains the SI-enhanced compressed latent to remain close to the target VAE latent:

$$\mathcal{L}_{guide} = \|\tilde{z}_{Dx} - x_0\|_2^2. \quad (46)$$

To supervise the encoder-side importance map, we derive a spatial reference-confidence map from the MCA scores:

$$A(i, j) = \max_k M_{Soft}^{(k)}(i, j). \quad (47)$$

This confidence measures the preference among retrieved references. It does not independently measure ab-

solute reference relevance and equals one when $K = 1$. The mask alignment loss is then defined as

$$\mathcal{L}_{IM} = \|M_{pred} - \text{sg}(A)\|_2^2, \quad (48)$$

where $\text{sg}(\cdot)$ denotes the stop-gradient operation.

V. EXPERIMENTAL RESULTS

A. Experimental Settings

1) *Datasets*. RADIC is trained on KITTI General and evaluated on KITTI General [10], KITTI Stereo [29], Cityscapes [30], and InStereo2K [31]. KITTI General contains 174936 training pairs, 912 validation pairs, and 3607 test pairs. The test sets contain 790 KITTI Stereo pairs, 1525 Cityscapes pairs, and 50 InStereo2K pairs.

2) *Decoder-Side SI Settings*. A fixed decoder-side gallery containing 23646 unique images is used during training. It includes 8147 images from KITTI General, 8389 from KITTI Stereo, 5000 from Cityscapes, and 2110 from InStereo2K. For the test comparison of paired, non-paired, and cross-domain SI, we use a test gallery containing 2369 images from the four datasets, including 394 KITTI General, 400 KITTI Stereo, 1525 Cityscapes, and 50 InStereo2K images.

The decoded target latent is used to query the corresponding FAISS index. During training, the number of SI references is sampled independently for each input as $K \sim \text{Uniform}\{0, 1, 2, 3\}$. For $K > 0$, FAISS retrieves the top-20 most similar gallery images, from which K references are randomly sampled.

For testing, paired SI retains paired or same-scene references. Non-paired SI removes the paired reference and all references from the same scene or sequence. Cross-domain SI removes references from the target domain. KITTI General and KITTI Stereo are treated as the same KITTI domain.

3) *Implementation Details*. All images are center-cropped to 370×740 and resized to 256×512 . Depth Anything V3 [28] is used for monocular depth estimation, while FAISS [32] is used for decoder-side retrieval. The diffusion decoder follows RDEIC with two denoising steps.

RADIC is trained using AdamW with an initial learning rate of 1×10^{-4} , and a weight decay of 1×10^{-2} . Cosine annealing is applied with a minimum learning rate of 1×10^{-6} . Each operating point is trained for 15 epochs with a batch size of 8 and a fixed random seed of 231. The loss coefficients are set to $\lambda_1 = 2.0$, $\lambda_3 = 2.0$, and $\lambda_4 = 0.1$, while $\lambda_2 \in \{16, 8, 4, 2, 1\}$ is varied to obtain different rate points. All experiments are conducted on NVIDIA L40 GPUs.

4) *Evaluation Metrics*. We use LPIPS [33], DISTS [34], FID [35], KID [36], and NIQE [37] for perceptual quality, and PSNR [38] and MS-SSIM [39] for pixel-level and structural fidelity. Perceptual metrics are used as the primary criteria at extremely low bitrates [11].




Fig. 6: Comparison of BD-Rate (%) [40], [41] computed with PSNR and MS-SSIM as quality metrics on four test datasets. RADIC is used as the anchor, and lower BD-Rate indicates better rate-distortion performance.

5) *Baselines*. We compare RADIC with the DIC methods NDIC [5], ATN [10], NDSC [6], LDMIC [9], MDIC [7], and CODIC [26]. We also include the joint multi-view codecs ECSIC [14], BiSIC [15], CAMSIC [16], and LMVIC [17], together with the single-image perceptual codecs RDEIC [13] and ResULIC [24]. All baselines are retrained on KITTI General.

B. Main Results

1) *Quantitative Comparison*. Fig. 5 compares perceptual performance at extremely low bitrates. RADIC achieves the best or near-best results on most perceptual metrics and maintains this advantage on KITTI Stereo, Cityscapes, and InStereo2K without dataset-specific fine-tuning. Under the common KITTI General training protocol, these results show favorable fixed-model transfer.

The perceptual gains do not consistently yield higher PSNR, reflecting the perception-distortion trade-off [42], [43]. As shown in Fig. 6, RADIC remains below several distortion-oriented methods in PSNR, particularly at higher rates, but achieves the best MS-SSIM on KITTI General and KITTI Stereo and remains competitive on Cityscapes and InStereo2K. This indicates that the perceptual improvements largely preserve structural fidelity despite lower pixel accuracy.

Compared with MDIC, RADIC replaces fixed paired SI with retrieved references. Its consistent gains across additional test datasets indicate reduced dependence on domain-specific paired SI. Compared with RDEIC and




Fig. 7: Qualitative comparison among CAMSIC [16], NDSC [6], MDIC [7], ResULIC [24], and RADIC. CAMSIC and NDSC exhibit severe degradation at extremely low bitrates, while RADIC has the lowest LPIPS in the displayed examples. PSNR varies by image and method.

TABLE I

BD-PSNR AND BD-LPIPS UNDER PAIRED, NON-PAIRED, AND CROSS-DOMAIN RETRIEVAL. RADIC WITH PAIRED SI AT $K = 1$ SERVES AS THE COMMON ANCHOR ACROSS DATASETS. HIGHER BD-PSNR AND LOWER BD-LPIPS ARE BETTER, WHILE BOLD DENOTES THE BEST RESULT IN EACH COLUMN.

Domain	Method	K	Paired		Non-paired		Cross-domain	
			BD-PSNR $\uparrow$	BD-LPIPS $\downarrow$	BD-PSNR $\uparrow$	BD-LPIPS $\downarrow$	BD-PSNR $\uparrow$	BD-LPIPS $\downarrow$
KITTI General	NDSC	1	1.26	0.1533	-1.19	0.2985	-1.23	0.3133
	ATN	1	0.13	0.2824	-0.29	0.3096	-0.34	0.3135
	MDIC	1	-2.49	0.0665	-4.27	0.1675	-4.32	0.1807
	RADIC	1	0.00	0.0000	-1.03	0.0233	-1.12	0.0471
		2	0.06	-0.0031	-1.03	0.0231	-1.12	0.0428
		3	0.08	-0.0043	-1.03	0.0233	-1.13	0.0401
KITTI Stereo	NDSC	1	0.27	0.2121	-1.09	0.2899	-2.35	0.3608
	ATN	1	-0.48	0.3076	-0.66	0.3180	-1.17	0.3419
	MDIC	1	-3.61	0.1112	-4.70	0.1638	-5.48	0.2259
	RADIC	1	0.00	0.0000	-0.86	0.0515	-0.95	0.0563
		2	0.08	-0.0014	-0.87	0.0512	-0.95	0.0561
		3	0.17	-0.0020	-0.86	0.0512	-0.95	0.0562
Cityscapes	NDSC	1	2.10	0.1458	-0.98	0.2888	-2.18	0.3660
	ATN	1	-0.42	0.3475	-0.64	0.3522	-0.92	0.3880
	MDIC	1	-3.24	0.0584	-5.44	0.1541	-6.09	0.2023
	RADIC	1	0.00	0.0000	-2.39	0.1291	-2.45	0.1375
		2	-0.16	0.0104	-2.38	0.1288	-2.44	0.1369
		3	-0.23	0.0135	-2.39	0.1304	-2.43	0.1371
InStereo2K	NDSC	1	1.05	0.1637	-1.70	0.2619	-2.22	0.2680
	ATN	1	0.77	0.2169	-0.21	0.2544	-0.35	0.2596
	MDIC	1	-1.94	0.0183	-3.99	0.1087	-4.75	0.1194
	RADIC	1	0.00	0.0000	-1.61	0.0812	-1.66	0.0816
		2	-0.18	-0.0074	-1.61	0.0808	-1.66	0.0819
		3	-0.28	0.0137	-1.61	0.0814	-1.67	0.0833

ResULIC, SI provides structural and textural information unavailable from the compressed target alone.

2) *Qualitative Comparison.* As shown in Fig. 7, NDSC and CAMSIC exhibit severe structural distortion and texture loss at extremely low bitrates. MDIC and ResULIC recover more realistic content but introduce noticeable deviations from the original image. In contrast, RADIC preserves clearer structures and finer textures at a lower bitrate, while maintaining higher pixel-level consistency with the target image.

3) *Robustness to SI Correlation.* Table I compares the

robustness of different DIC methods under paired, non-paired, and cross-domain SI. For a fair comparison, each baseline uses the same top-1 retrieved SI as RADIC.

RADIC achieves the lowest BD-LPIPS across all datasets and retrieval settings, showing that its perceptual reconstruction remains stable as SI correlation decreases. The transition from paired to cross-domain SI causes substantially smaller degradation than MDIC in both BD-LPIPS and BD-PSNR across the four datasets. This indicates that DAGA and MCA reduce the effect of geometric mismatch and irrelevant content. ATN is gen-

TABLE II
BD-RATE (%) ↓ OF RADIC VARIANTS COMPUTED USING LPIPS, DISTIS, PSNR, AND MS-SSIM. THE FULL RADIC MODEL IS USED AS THE ANCHOR.

Dataset	Variant	Components				BD-Rate (%) ↓			
		Ref.	DAGA	MCA	IM	LPIPS	DISTS	PSNR	MS-SSIM
KITTI General	C1	–	–	–	–	166.8	134.3	163.5	175.4
	C2	✓	–	–	–	81.0	65.8	87.2	88.5
	C3	✓	✓	–	–	13.8	11.5	14.9	14.9
	C4	✓	✓	✓	–	1.3	-1.8	11.3	1.8
	RADIC	✓	✓	✓	✓	0.0	0.0	0.0	0.0
KITTI Stereo	C1	–	–	–	–	182.1	127.5	189.7	197.9
	C2	✓	–	–	–	95.4	81.4	95.0	97.0
	C3	✓	✓	–	–	15.7	13.4	18.0	16.2
	C4	✓	✓	✓	–	11.3	12.5	11.2	13.9
	RADIC	✓	✓	✓	✓	0.0	0.0	0.0	0.0
Cityscapes	C1	–	–	–	–	38.7	38.6	67.7	115.6
	C2	✓	–	–	–	36.6	35.0	58.7	110.8
	C3	✓	✓	–	–	27.3	27.6	50.1	65.8
	C4	✓	✓	✓	–	1.1	0.7	31.5	0.5
	RADIC	✓	✓	✓	✓	0.0	0.0	0.0	0.0
InStereo2K	C1	–	–	–	–	132.8	131.7	94.8	172.2
	C2	✓	–	–	–	38.8	30.9	45.2	81.2
	C3	✓	✓	–	–	10.6	15.9	18.9	14.5
	C4	✓	✓	✓	–	6.6	4.0	10.7	10.9
	RADIC	✓	✓	✓	✓	0.0	0.0	0.0	0.0

erally more robust than NDSC in BD-PSNR and in several cross-domain perceptual comparisons, which is consistent with its use of cross-attention for feature selection, although it remains notably below RADIC in BD-LPIPS. Increasing K from 1 to 3 produces only marginal changes under non-paired and cross-domain retrieval, indicating that the first retrieved reference provides most of the available information and that additional weakly related candidates are largely suppressed. In this paper, $K = 3$ is used to balance reference coverage and interference for all main results.

4) *Complexity Analysis.* As shown in Table III, RADIC matches RDEIC with the lowest encoding latency of 0.0149 s. Despite the added cost of retrieval and alignment, its decoding remains faster than MDIC and ResULIC. Using RADIC as the anchor, MDIC requires 17.68% and 136.11% higher rates for LPIPS and PSNR, respectively, while the single-image codecs show larger gaps.

C. Ablation Studies

Table II reports the cumulative ablation results of RADIC. C_1 denotes the base codec. Reference retrieval, DAGA, and MCA are introduced in C_2 , C_3 , and C_4 , while the full model further includes the importance map. 1) *Reference Retrieval.* Adding retrieved SI from C_1 to C_2 reduces BD-Rate for all four metrics on all datasets. The improvements are substantial on KITTI General, KITTI Stereo, and InStereo2K. The gain is smaller

TABLE III
RUNTIME AND BD-RATE COMPARISON OF PERCEPTION-ORIENTED METHODS ON CITYSCAPES. ANCHOR: RADIC ($K = 3$).

Model	Time (s)		BD-Rate (%) ↓	
	Enc.	Dec.	LPIPS	PSNR
ResULIC	0.0749 ± 0.0136	0.5143 ± 0.0158	151.02	238.63
RDEIC	0.0149 ± 0.0001	0.0689 ± 0.0001	81.32	112.59
MDIC	0.0331 ± 0.0026	0.4244 ± 0.0128	17.68	136.11
RADIC	0.0149 ± 0.0001	0.2357 ± 0.0016	0.00	0.00

on Cityscapes, indicating that the benefit of retrieval varies with the available SI. These results show that retrieved references provide useful information beyond the compressed target.

2) *Depth-Aware Geometric Alignment.* Introducing DAGA in C_3 further reduces BD-Rate across all datasets and metrics. The reductions are particularly large on KITTI General, KITTI Stereo, and InStereo2K. Cityscapes also shows clear improvement, especially in MS-SSIM. The consistent gains show the importance of reducing geometric differences between the target and retrieved references before feature fusion.

3) *Multi-reference Cross-Attention.* Adding MCA in C_4 provides further gains after geometric alignment. The largest improvements are observed on Cityscapes and KITTI General, while the changes on KITTI Stereo and InStereo2K are smaller. This result is consistent with the




Fig. 8: Effect of the number of SI references under paired, non-paired, and cross-domain SI. Rate saving is defined as the negative BD-Rate relative to $K = 0$ under the same SI setting, where higher values indicate better coding efficiency.

role of MCA in selecting and combining information from the aligned references.

4) *SI-Aware Rate Allocation*. The full RADIC model improves most metric combinations over C_4 . Clear gains are observed across all four metrics on KITTI Stereo and InStereo2K, while the PSNR improvement is especially large on Cityscapes. On KITTI General, C_4 achieves a slightly lower BD-Rate for DISTS, with a difference of 1.8%. The overall results show that the importance map provides additional gains through adaptive modulation of the transmitted latent.

VI. DISCUSSION

The ablation results show that retrieved SI improves compression performance, and that DAGA and MCA provide further gains after retrieval. We next examine how the number of references affects the results under different SI conditions.

Fig. 8 reports the rate saving from $K = 0$ to $K = 6$ under paired, non-paired, and cross-domain SI. Rate saving is defined as the negative BD-Rate relative to $K = 0$ under the same SI setting. PSNR provides the closest empirical comparison with the distortion term in the information-theoretic analysis because it is directly related to MSE for a fixed image range.

Under paired SI, the first reference accounts for most of the rate saving on all four datasets. Increasing K beyond one produces only small changes. KITTI General

shows slight additional gains for several metrics, while the other datasets mainly show saturation or gradual decreases. This observation is consistent with the condition

$$I(U; Y_{K+1} | Y_{1:K}) > 0, \quad (49)$$

which describes whether an additional reference contains information not already provided by the existing SI.

The gains decrease when the available SI has lower correlation with the target. Under non-paired SI, KITTI General, KITTI Stereo, and InStereo2K show mostly negative rate savings. Cityscapes retains small gains for some metrics. Under cross-domain SI, the rate savings are negative across the four datasets and metrics. The ablation results show that DAGA and MCA improve the use of retrieved references, but it does not compensate for low reference relevance.

PSNR, LPIPS, and DISTS do not always follow the same trend. LPIPS and DISTS evaluate perceptual similarity in feature space, while PSNR is derived from pixel-level distortion. The main experiments use $K = 3$, which is also the largest number of references used during training. Larger values of K provide little additional gain.

VII. CONCLUSION

We present RADIC, a distributed image compression framework that retrieves multiple side information (SI) references from a decoder-side visual gallery for extremely low-bitrate reconstruction. Retrieved references

are aligned by the Depth-Aware Geometric Alignment (DAGA) module and fused by the Multiple Reference Cross-Attention (MCA) module to handle geometric and content differences. Decoder attention is also used to supervise an encoder-side importance map for SI-aware rate allocation, while the encoder uses only the target image. Experiments on KITTI General, KITTI Stereo, Cityscapes, and InStereo2K show that RADIC achieves strong perceptual compression performance and maintains stable reconstruction quality under paired, non-paired, and cross-domain SI. The analysis of the number of references shows that most of the gain is obtained from a small number of relevant references. Additional references provide limited benefit when their correlation with the target becomes weak. These results demonstrate the effectiveness of retrieved SI for distributed image compression and highlight the importance of reference relevance in decoder-side visual galleries.

Limitations. Although RADIC is more efficient than existing perceptual DIC methods, further reducing retrieval and diffusion costs remains important. As a new cloud-edge compression paradigm, broader exploration of distributed data in open scenarios is needed to support scalable deployment.

REFERENCES

- [1] C. Anthes, R. J. García-Hernández, M. Wiedemann, and D. Kranzlmüller, "State of the art of virtual reality technology," in *2016 IEEE Aerospace Conference*. IEEE, 2016, pp. 1–19.
- [2] C. Prenebida, R. Ambrus, and Z.-C. Marton, "Intelligent robotic perception systems," *Applications of Mobile Robots*, pp. 111–127, 2018.
- [3] C. Anthes, R. J. García-Hernández, M. Wiedemann, and D. Kranzlmüller, "State of the art of virtual reality technology," in *2016 IEEE Aerospace Conference*. IEEE, 2016, pp. 1–19.
- [4] P. T. Endo, A. V. de Almeida Palhares, N. N. Pereira, G. E. Gonçalves, D. Sadok, J. Kelner, B. Melander, and J.-E. Mangs, "Resource allocation for distributed cloud: concepts and research challenges," *IEEE Network*, vol. 25, no. 4, pp. 42–46, 2011.
- [5] N. Mital, E. Özyilkan, A. Garjani, and D. Gündüz, "Neural distributed image compression using common information," in *2022 Data Compression Conference (DCC)*. IEEE, 2022, pp. 182–191.
- [6] J. Whang, A. Nagle, A. Acharya, H. Kim, and A. G. Dimakis, "Neural distributed source coding," *IEEE Journal on Selected Areas in Information Theory*, vol. 5, pp. 493–508, 2024.
- [7] G. Xu, M. Zhang, J. Xiang, C. Tan, Y. Yang, and J. Zhou, "Distributed image compression with multimodal side information at extremely low bitrates," in *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, June 2026, pp. 33 662–33 671.
- [8] Y. Huang, B. Chen, S. Qin, J. Li, Y. Wang, T. Dai, and S.-T. Xia, "Learned distributed image compression with multi-scale patch matching in feature domain," in *Proceedings of the AAAI Conference on Artificial Intelligence*, vol. 37, no. 4, 2023, pp. 4322–4329.
- [9] X. Zhang, J. Shao, and J. Zhang, "Ldmic: Learning-based distributed multi-view image coding," in *International Conference on Learning Representations*, 2023.
- [10] N. Mital, E. Özyilkan, A. Garjani, and D. Gündüz, "Neural distributed image compression with cross-attention feature alignment," in *Proceedings of the IEEE/CVF Winter Conference on Applications of Computer Vision*, 2023, pp. 2498–2507.
- [11] M. Careil, M. J. Muckley, J. Verbeek, and S. Lathuilière, "Towards image compression with perfect realism at ultra-low bitrates," in *The Twelfth International Conference on Learning Representations*, 2023.
- [12] T. Bordin and T. Maugey, "Linearly transformed color guide for low-bitrate diffusion-based image compression," *IEEE Transactions on Image Processing*, vol. 34, pp. 468–482, 2025.
- [13] Z. Li, Y. Zhou, H. Wei, C. Ge, and A. Mian, "Rdeic: Accelerating diffusion-based extreme image compression with relay residual diffusion," *IEEE Transactions on Circuits and Systems for Video Technology*, pp. 1–1, 2025.
- [14] M. Wödlinger, J. Kotera, M. Keglevic, J. Xu, and R. Sablatnig, "Ecsic: Epipolar cross attention for stereo image compression," in *Proceedings of the IEEE/CVF Winter Conference on Applications of Computer Vision*, 2024, pp. 3436–3445.
- [15] Z. Liu, X. Zhang, J. Shao, Z. Lin, and J. Zhang, "Bidirectional stereo image compression with cross-dimensional entropy model," in *European Conference on Computer Vision*. Springer, 2024, pp. 480–496.
- [16] X. Zhang, S. Gao, Z. Liu, J. Shao, X. Ge, D. He, T. Xu, Y. Wang, and J. Zhang, "Camsic: Content-aware masked image modeling transformer for stereo image compression," in *Proceedings of the AAAI Conference on Artificial Intelligence*, vol. 39, no. 10, 2025, pp. 10 239–10 247.
- [17] Y. Huang, B. Chen, N. Lian, X. Wang, B. An, T. Dai, and S.-T. Xia, "3D-LMVIC: Learning-based multi-view image compression with 3D Gaussian geometric priors," in *Proceedings of the 42nd International Conference on Machine Learning*, ser. Proceedings of Machine Learning Research, A. Singh, M. Fazel, D. Hsu, S. Lacoste-Julien, F. Berkenkamp, T. Maharaj, K. Wagstaff, and J. Zhu, Eds., vol. 267. PMLR, 13–19 Jul 2025, pp. 25 141–25 158. [Online]. Available: <https://proceedings.mlr.press/v267/huang25e.html>
- [18] D. Slepian and J. Wolf, "Noiseless coding of correlated information sources," *IEEE Transactions on Information Theory*, vol. 19, no. 4, pp. 471–480, 1973.
- [19] A. Wyner and J. Ziv, "The rate-distortion function for source coding with side information at the decoder," *IEEE Transactions on Information Theory*, vol. 22, no. 1, pp. 1–10, 2003.
- [20] H. Yue, X. Li, M. Sun, G. Liu, X. Liu, and J. Yang, "DIC-DDA: Learned asymmetric distributed image compression via dual domain alignment," *IEEE Transactions on Image Processing*, vol. 35, pp. 3822–3834, 2026.
- [21] Z. Li, Y. Zhou, H. Wei, C. Ge, and J. Jiang, "Towards extreme image compression with latent feature guidance and diffusion prior," *IEEE Transactions on Circuits and Systems for Video Technology*, 2024.
- [22] Q. Mao, C. Wang, M. Wang, S. Wang, R. Chen, L. Jin, and S. Ma, "Scalable face image coding via StyleGAN prior: Toward compression for human-machine collaborative vision," *IEEE Transactions on Image Processing*, vol. 33, pp. 408–422, 2024.
- [23] J. Chang, Z. Zhao, C. Jia, S. Wang, L. Yang, Q. Mao, J. Zhang, and S. Ma, "Conceptual compression via deep structure and texture synthesis," *IEEE Transactions on Image Processing*, vol. 31, pp. 2809–2823, 2022.
- [24] A. Ke, X. Zhang, T. Chen, M. Lu, C. Zhou, J. Gu, and Z. Ma, "Ultra lowrate image compression with semantic residual coding and compression-aware diffusion," in *International Conference on Machine Learning*, 2025.
- [25] J. Gao, Z. Huang, Q. Mao, S. Ma, and C. Jia, "Exploring multimodal knowledge for image compression via large foundation models," *IEEE Transactions on Image Processing*, vol. 34, pp. 5904–5919, 2025.
- [26] G. Xu, C. Tan, J. Xiang, and J. Zhou, "CODIC: Distributed image compression with distortion-perception co-optimization," *IEEE Transactions on Multimedia*, 2026, early Access.
- [27] A. Van Den Oord, O. Vinyals *et al.*, "Neural discrete representation learning," *Advances in Neural Information Processing Systems*, vol. 30, 2017.
- [28] H. Lin, S. Chen, J. Liew, D. Y. Chen, Z. Li, G. Shi, J. Feng, and B. Kang, "Depth anything 3: Recovering the visual space from any views," *arXiv preprint arXiv:2511.10647*, 2025.

- [29] A. Geiger, P. Lenz, and R. Urtasun, "Are we ready for autonomous driving? the kitti vision benchmark suite," in *2012 IEEE conference on computer vision and pattern recognition*. IEEE, 2012, pp. 3354–3361.
- [30] M. Cordts, M. Omran, S. Ramos, T. Rehfeld, M. Enzweiler, R. Benenson, U. Franke, S. Roth, and B. Schiele, "The cityscapes dataset for semantic urban scene understanding," in *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition*, 2016, pp. 3213–3223.
- [31] W. Bao, W. Wang, Y. Xu, Y. Guo, S. Hong, and X. Zhang, "Instereo2k: a large real dataset for stereo matching in indoor scenes," *Science China Information Sciences*, vol. 63, no. 11, p. 212101, 2020.
- [32] M. Douze, A. Guzhva, C. Deng, J. Johnson, G. Szilvasy, P.-E. Mazaré, M. Lomeli, L. Hosseini, and H. Jégou, "The faiss library," *IEEE Transactions on Big Data*, 2025.
- [33] R. Zhang, P. Isola, A. A. Efros, E. Shechtman, and O. Wang, "The unreasonable effectiveness of deep features as a perceptual metric," in *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition*, 2018, pp. 586–595.
- [34] K. Ding, K. Ma, S. Wang, and E. P. Simoncelli, "Image quality assessment: Unifying structure and texture similarity," *IEEE Transactions on Pattern Analysis and Machine Intelligence*, vol. 44, no. 5, pp. 2567–2581, 2020.
- [35] M. Heusel, H. Ramsauer, T. Unterthiner, B. Nessler, and S. Hochreiter, "Gans trained by a two time-scale update rule converge to a local nash equilibrium," *Advances in Neural Information Processing Systems*, vol. 30, 2017.
- [36] M. Bińkowski, D. J. Sutherland, M. Arbel, and A. Gretton, "Demystifying mmd gans," *arXiv preprint arXiv:1801.01401*, 2018.
- [37] A. Mittal, R. Soundararajan, and A. C. Bovik, "Making a "completely blind" image quality analyzer," *IEEE Signal Processing Letters*, vol. 20, no. 3, pp. 209–212, 2012.
- [38] J. Korhonen and J. You, "Peak signal-to-noise ratio revisited: Is simple beautiful?" in *2012 Fourth International Workshop on Quality of Multimedia Experience*, 2012, pp. 37–38.
- [39] Z. Wang, E. P. Simoncelli, and A. C. Bovik, "Multiscale structural similarity for image quality assessment," in *The Thirty-seventh Asilomar Conference on Signals, Systems & Computers*, 2003, vol. 2. IEEE, 2003, pp. 1398–1402.
- [40] G. Bjontegaard, "Calculation of average psnr differences between rd-curves," *ITU SG16 Doc. VCEG-M33*, 2001.
- [41] C. Herglotz, H. Och, A. Meyer, G. Ramasubbu, L. Eicher Müller, M. Kränzler, F. Brand, K. Fischer, D. T. Nguyen, A. Regensky *et al.*, "The Bjontegaard bible: Why your way of comparing video codecs may be wrong," *IEEE Transactions on Image Processing*, vol. 33, pp. 987–1001, 2024.
- [42] C. Zhou, G. Lu, J. Li, X. Chen, Z. Cheng, L. Song, and W. Zhang, "Controllable distortion-perception tradeoff through latent diffusion for neural image compression," in *Proceedings of the AAAI Conference on Artificial Intelligence*, vol. 39, no. 10, 2025, pp. 10 725–10 733.
- [43] Y. Blau and T. Michaeli, "Rethinking lossy compression: The rate-distortion-perception tradeoff," in *International Conference on Machine Learning*. PMLR, 2019, pp. 675–685.



Guojun Xu is currently pursuing a Ph.D. degree with the School of Artificial Intelligence, Wuhan University of Technology. His current research interests include distributed source coding, image compression, and other visual problems.



Jianwen Xiang (Member, IEEE) received the Ph.D. degree in computer software and theory from Wuhan University and in information science from the Japan Advanced Institute of Science and Technology. Since then, he has been a Special Research Fellow with the National Institute of Advanced Industrial Science and Technology, Japan, the Director of the Central Research Laboratory, NEC Corporation, Japan, and a Postdoctoral Researcher with the Japan Advanced Institute of Science and Technology. He is currently a Professor with the School of Computer Science and Technology, Wuhan University of Technology, China, and a Visiting Professor with Kyushu University, Japan. His research interests include dependable computing, information security, and formal methods. He has served as a Reviewer for international journals, including IEEE Transactions on Reliability, Reliability Engineering and System Safety, ACM Transactions on Software Engineering and Methodology, and Journal of Systems and Software.



Mingyang Zhang is currently pursuing the M.S. degree with the School of Artificial Intelligence, Wuhan University of Technology. His current research interests include distributed source coding and multi-reference face super-resolution.

Yanling Xie is currently pursuing the M.S. degree with the School of Artificial Intelligence, Wuhan University of Technology. His current research interests include distributed source coding, lossless image compression, and computer vision.



Junwei Zhou (Member, IEEE) received the Ph.D. degree from the Department of Electrical Engineering, City University of Hong Kong. Since then, he has been a Visiting Scholar with the Polytechnic University of Turin, Italy, and a Researcher with the City University of Hong Kong, Hong Kong, China. He is currently a Professor with the School of Artificial Intelligence, Wuhan University of Technology, China, and a Postdoctoral Researcher with Pennsylvania State University, USA. His research interests include computer vision, information security, and distributed source coding. He has served as a PC Member for conferences including EMNLP 2023, Coling 2022, ACL 2022, and ISSRE 2022.